

# Unveiling the Power of Deep Learning

Autoencoders and Generative Adversarial Networks (GANs)

Dr. Muhammad Sajjad

R.A: Kaleem Ullah

R.A: Aizaz Ali Shah

# Overview

## ➤ **Autoencoders:**

- Concept and Components
- Properties of autoencoder
- Types of Autoencoders
  - Vanilla AutoEncoders
  - Convolutional AutoEncoders
  - Denoising AutoEncoders
- Applications of Autoencoders

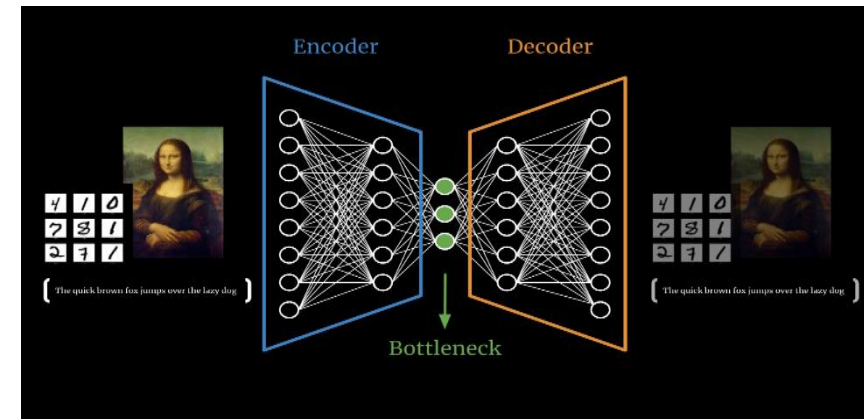

## ➤ **Generative Adversarial Networks (GANs):**

- Concept and Components
- Types of GANs
  - DCGANs
  - CGANs
  - StackGANs
- Applications of GANs

# What are autoencoder?

## Autoencoder:

- Autoencoders are a type of artificial neural network that aim to learn efficient representations (codings, embeddings) of input data.
- They are designed to encode the input into a compressed and reduced representation and then decode this representation back to the original like or any other desired output.
- This helps the network learn important features and patterns within the data.
- They are trained to learn efficient representations of input data, like images, text, etc.



# Architecture of Autoencoders

**Components:** Composed of Encoder, Decoder And Latent Space

## 1. Encoder:

The encoder compresses the input data into a latent space representation.

It reduces the dimensionality of the input by learning the most salient features.

Mathematically, if  $\mathbf{x}$  is the input and  $\mathbf{h}$  is the latent space representation, the encoder function can be represented as  $\mathbf{h} = f(\mathbf{x})$ .

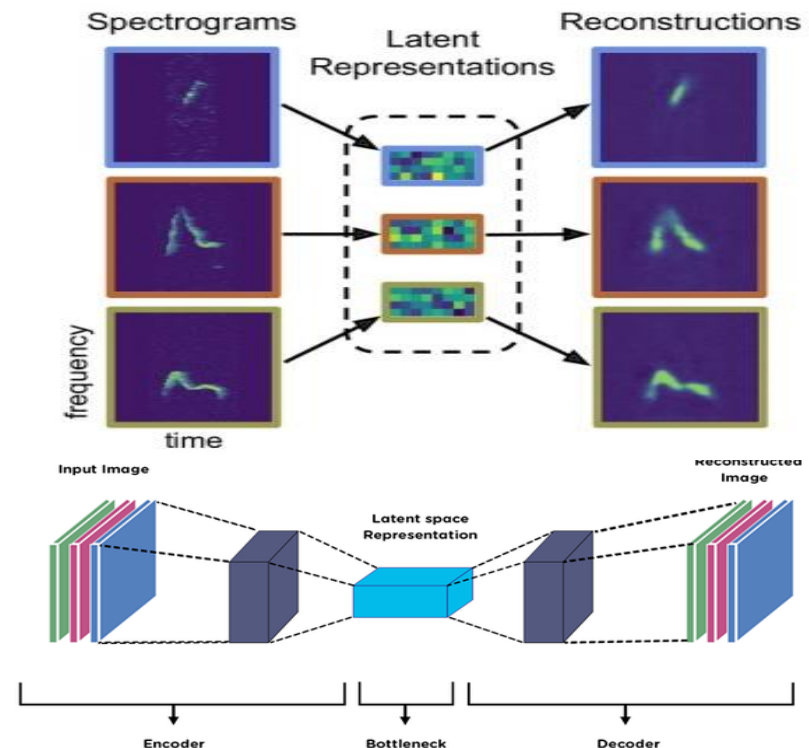

## 2. Latent Space:

The compressed, encoded representation of the input data.

Typically has a lower dimensionality than the input data.



## 3. Decoder:

- The decoder reconstructs the input data from the latent space representation.
- It attempts to recreate the input data as accurately as possible.
- Mathematically, if  $\hat{\mathbf{x}}$  is the reconstructed input, the decoder function can be represented as  $\hat{\mathbf{x}} = g(\mathbf{h})$ . Where  $\mathbf{h}$  is the encoded latent representation



# Feature Extraction and Dimensionality Reduction

- Compresses the facial image from 2000 dimensions to 30 dimensions.
- Utilizes multiple layers to achieve dimensionality reduction.
- Reconstructs the input image from the compressed representation.
- Ensures the autoencoder captures essential features during training.
- This approach highlights the utility of neural networks in feature extraction and image analysis.



# Properties of Autoencoder

## 1. Data-Specific

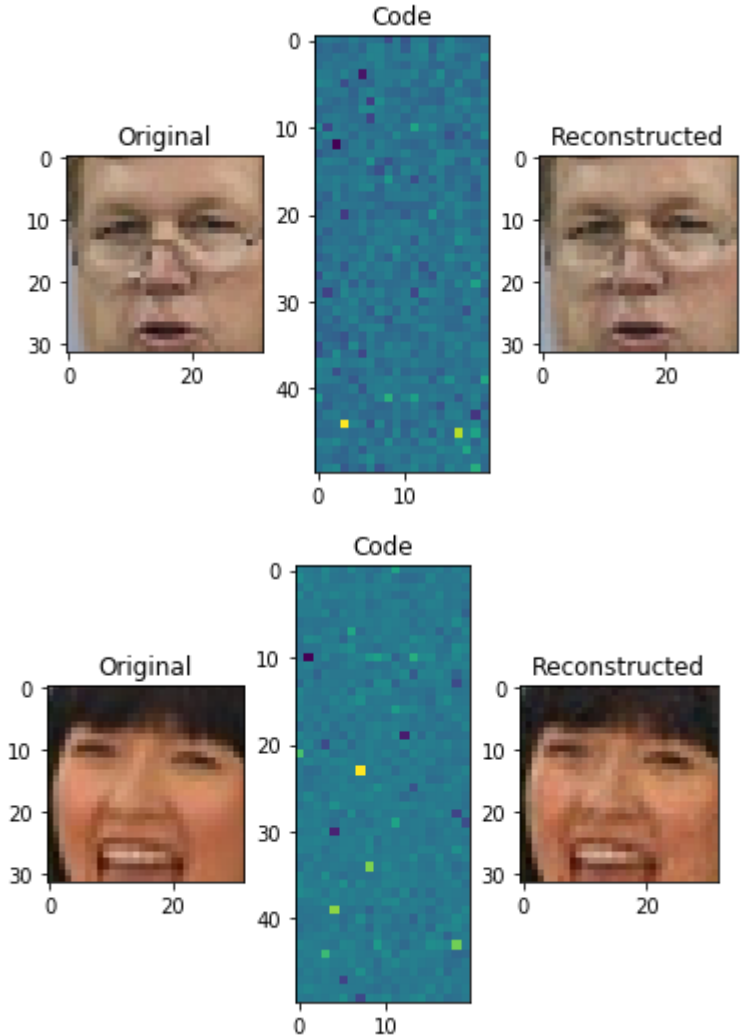

- Autoencoders are trained on a specific dataset, encoding and decoding the unique characteristics of that data

### Implications:

**Customization:** Autoencoders must be retrained for each new dataset or type of data you wish to process.


**Specialization:** The learned representations are optimized for the specific patterns and structures present in the training data.

**Generalization:** They may not generalize well to data significantly different from the training set without further training or adaptation.



## 2. Lossy

- Autoencoders are lossy compression methods, meaning that the reconstruction of the original data from the compressed representation is not perfect. There is some loss of information, which can be controlled and minimized but not entirely eliminated.




### Implications:

**Approximation:** The reconstructed data is an approximation of the original input, with some details potentially lost or altered.

**Compression Trade-off:** There is a trade-off between the level of compression and the accuracy of reconstruction. Higher compression typically leads to greater loss of detail.



**Application Suitability:** They are suitable for applications where perfect reconstruction is not critical, but where capturing the essential features of the data is sufficient (e.g., image compression, noise reduction).



# Types of autoencoder

- **Vanilla Autoencoder**



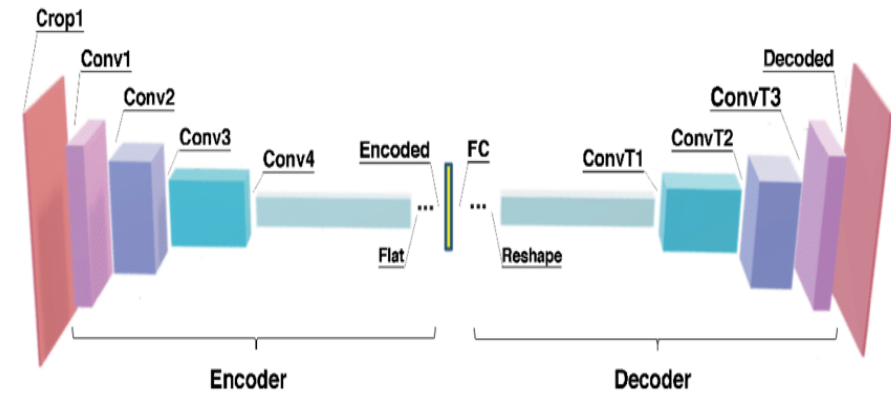
- A basic autoencoder with an encoder network compressing input data and a decoder network reconstructing it.
- **Purpose:**
  - Learn to compress input data into a low-dimensional representation.
  - Recover original input from this representation.
  - Preserve crucial aspects of the input data while minimizing data loss during compression.
- **Architecture:**
  - Encoder: Converts input data to a lower-dimensional representation.
  - Decoder: Converts this representation back to the original input.
- **Working:**
  - Minimizes reconstruction error between the original input and decoder output.
  - Uses loss functions like Mean Squared Error (MSE) or Binary Cross-Entropy (BCE).





# Convolutional Autoencoder


- An autoencoder designed specifically for image data, utilizing convolutional layers in the encoder and decoder networks.
- **Purpose:**
  - Discover a compressed representation of an input image.
  - Preserve spatial information.
- **Architecture:**
  - Uses convolutional layers in both encoder and decoder networks.
  - Extracts and preserves spatial information from input images.
- **Working:**
  - Minimizes reconstruction error between the original input image and the decoder output.




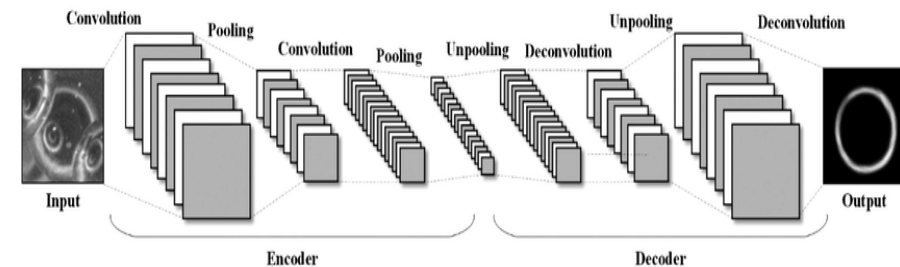


## • Denoising Autoencoder

- Trained with corrupted input data by adding noise and clean data as the objective, autoencoders aim to remove noise from the input data
- **Purpose:**
  - Learn to filter out noise from input data.
  - Develop a robust representation of data that is less sensitive to noise.
- **Working:**
  - Minimizes reconstruction error between clean input data and decoder output from corrupted data.
- **Applications:**
  - Voice and image denoising

### Denoising Autoencoders







### Visualization



# Application of Autoencoders

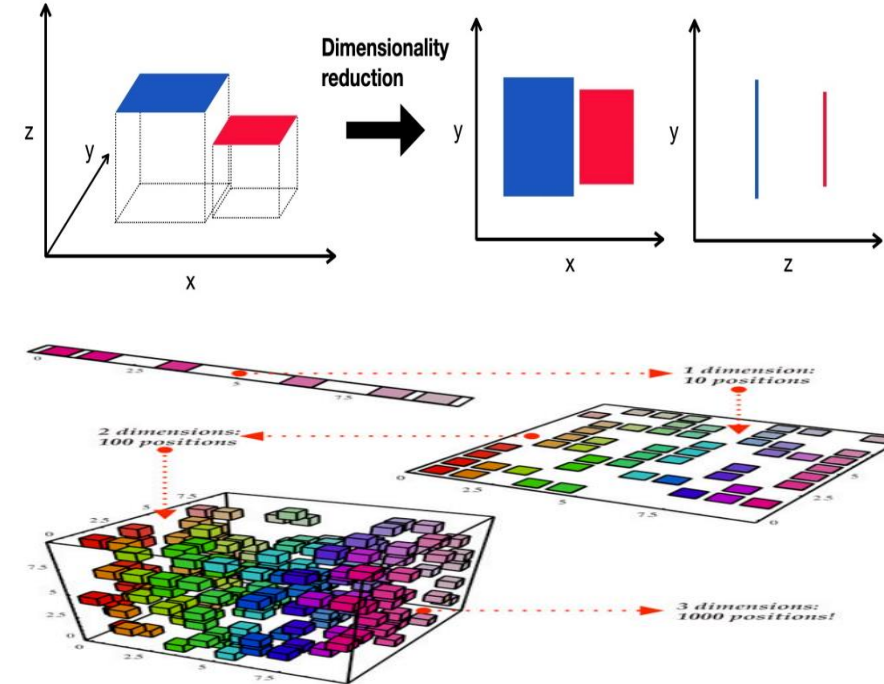

## 1. Data Denoising

- Removing noise from images, audio signals, or other data.
- **Image Denoising:**
  - Autoencoders reconstruct images while removing noise.
  - Compress input image into a latent space and reconstruct it.
  - Model learns to ignore noise and retain important features.
- **Audio Denoising:**
  - Clean noisy audio signals.
  - Enhances clarity and understandability.
- **Real-Time Example:**
  - Medical Imaging: Improves MRI or CT scan quality by removing noise, aiding in better diagnosis.



## 2. Dimensionality Reduction


- Reducing the number of features in a dataset while preserving important information.
- Compress high-dimensional data into a lower-dimensional latent space.
- Useful for visualization, reducing computational costs, and mitigating the curse of dimensionality.
- **Real-Time Example:**
  - Feature Extraction for Machine Learning:
    - In fields like genomics or image recognition, autoencoders reduce feature numbers, enhancing model speed and efficiency.
- **Data Visualization:**
  - Converts high-dimensional data (e.g., word embeddings, sensor data) to 2D or 3D for easier visualization and interpretation.







### 3. Image Compression

- Reducing the size of images for storage or transmission without significant loss of quality.
- Autoencoders learn to compress images into a smaller size and then reconstruct them.
- Achieves better compression rates and quality retention compared to traditional methods (e.g., JPEG) by learning sophisticated representations.
- **Real-Time Example:**
- **Web and Mobile Applications:**
- Reducing image sizes speeds up loading times and reduces bandwidth usage, enhancing user experience.



## 4. Semantic Segmentation

- Dividing an image into segments based on the content and context of different regions.
- Autoencoders encode complex images into simpler, more interpretable representations.
- Facilitates segmenting and labeling different parts of an image.
- **Real-Time Example:**
- **Autonomous Driving:**
- Recognizes and segments parts of the road scene (e.g., vehicles, pedestrians, traffic signs).
- Enhances the safety and efficiency of autonomous vehicles.



## 5. Anomaly Detection

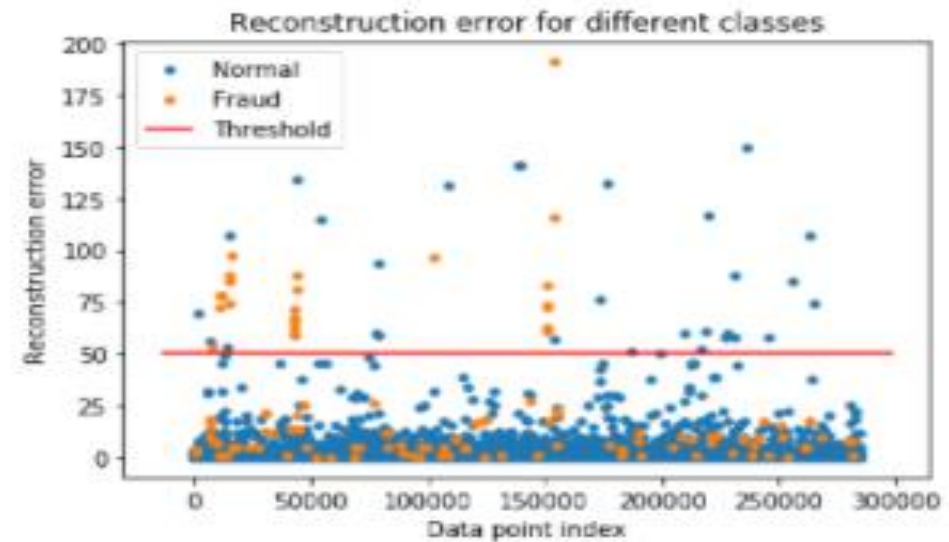
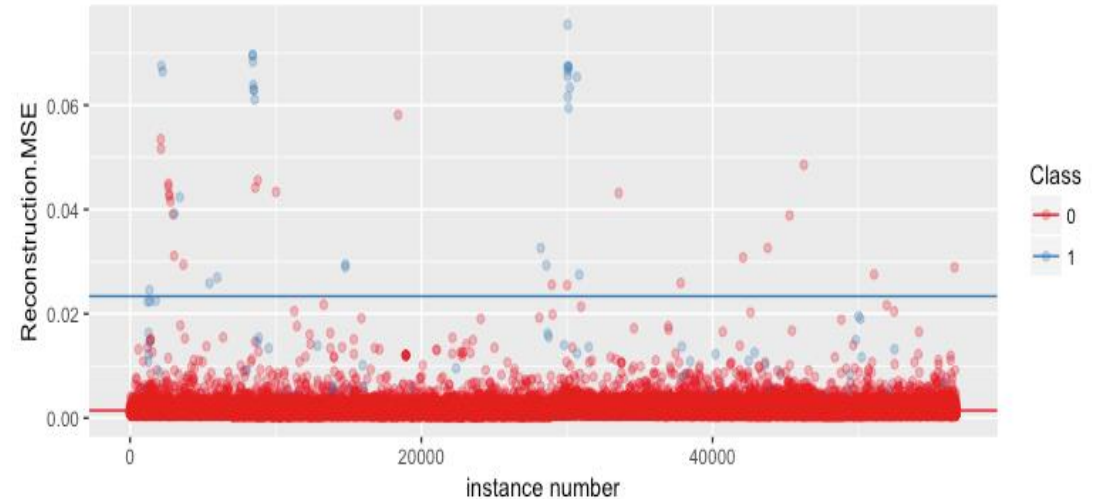
- **Real-Life Applications of Autoencoders**

- **Fraud Detection:**

- Detecting fraudulent transactions in banks.
- Trained on normal transactions; high reconstruction error flags potential fraud.
- Improves security by identifying unusual patterns indicative of fraud.

- **Equipment Monitoring:**

- Monitoring factory machines for potential failures.
- Trained on normal machine operation data; spikes in reconstruction error indicate anomalies.
- Enhances maintenance by early detection of equipment malfunctions.



- **Challenges and Limitations of Autoencoders**

- **Inability to Generate New Data**

- **Issue:** Autoencoders are designed to reconstruct input data rather than generate new, unseen data.
- **Solution:** For generative tasks, consider using models specifically designed for generating new data, such as Generative Adversarial Networks (GANs).

- **Overfitting**

- **Issue:** Learns noise along with data patterns.
- **Solution:** Use regularization, dropout, and early stopping.

- **Scalability**

- **Issue:** Requires significant computational power and time for large datasets.
- **Solution:** Utilize hardware acceleration (GPUs/TPUs), batch training, and efficient algorithms.



# Unveiling the Power of Creation

## Generative Adversarial Networks (GANs)

- **Origin and Significance of GANs**

- **Origin:**


- **Ian Goodfellow and Team:** Conceptualized by Ian Goodfellow, a PhD student at the University of Montreal under Yoshua Bengio, alongside Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio.
- **Paper Presentation:** Introduced in the "Generative Adversarial Nets" paper at NeurIPS 2014, a premier machine learning conference.
- **Generative Adversarial Nets Paper:**
- **Novel Approach:** Utilized two neural networks (generator and discriminator) in competition to produce realistic synthetic data.
- **Impact:** Marked a significant improvement over previous generative models.

- **Significance:**

- **AI Breakthrough:** GANs revolutionized machine learning and AI by enabling the generation of highly realistic images and data through an adversarial process.
- **Adversarial Process:** Continuous improvement of the generator and discriminator networks enhances data quality, surpassing earlier probabilistic methods and neural network architectures.

# What is a GAN?



- A generative adversarial network, or GAN, is a deep neural network framework which is able to learn from a set of training data and generate new data with the same characteristics as the training data
- **Discriminative Model:**
  - The discriminator's role is to distinguish between real and fake data samples.
  - It takes an input sample, processes it through a series of neural network layers, and outputs a probability score indicating whether the sample is real or fake.
- **Generative Model:**
  - The generator's role is to create new data samples that look similar to the training data.
  - It starts with a random noise vector. This vector is transformed through several layers of the neural network, which include operations such as fully connected layers, convolutional layers, and activation functions.



# GAN architecture

GANs are based on the idea of adversarial training. The GAN architecture basically consists of two neural networks that compete against each other:

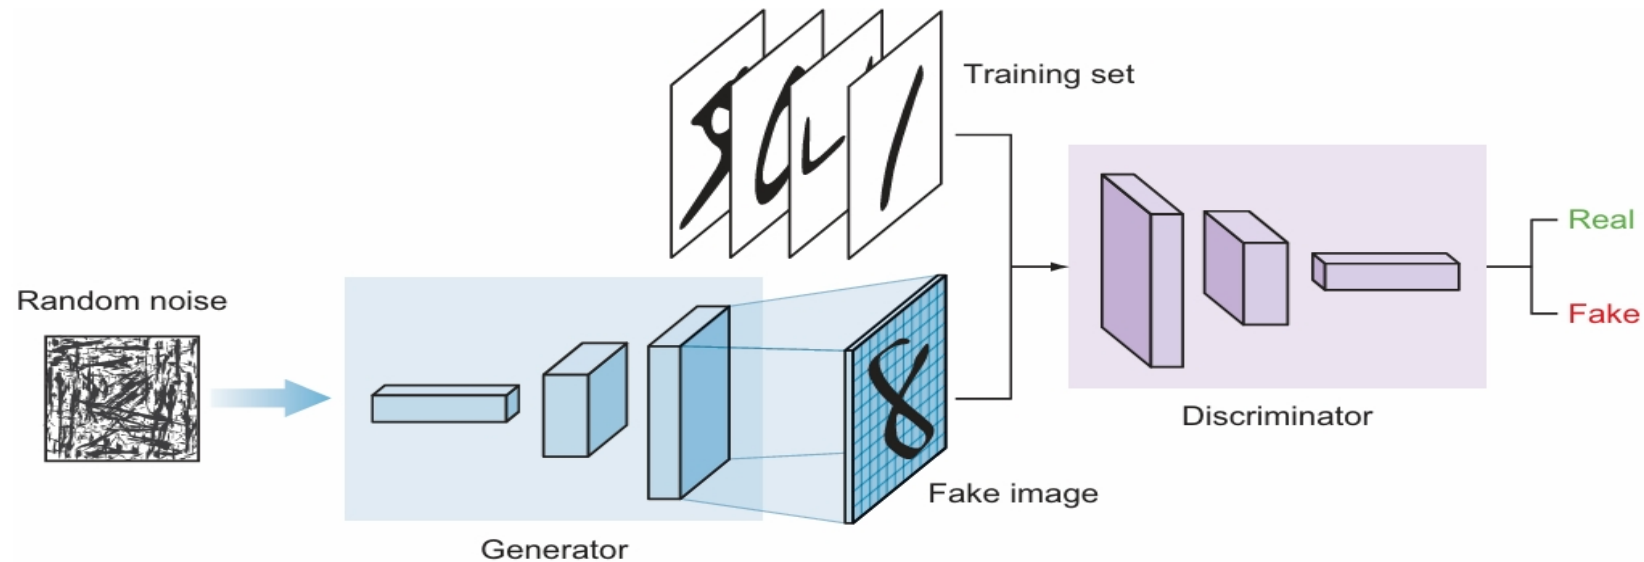
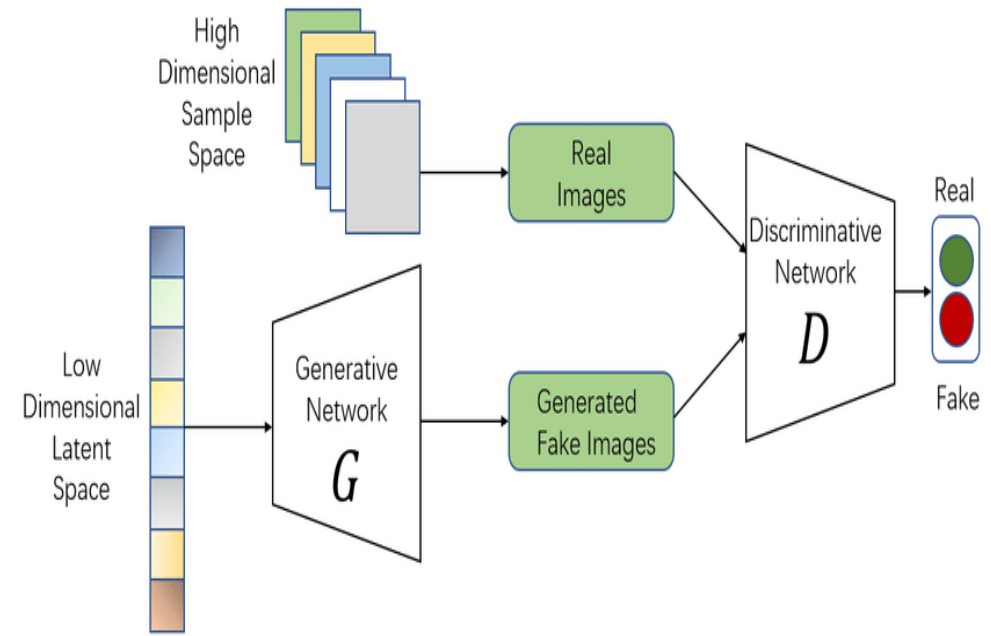
- The generator tries to convert random noise into images that look as if they have been sampled from the original dataset.
- The discriminator tries to predict whether an observation comes from the original dataset or is one of the generator's forgeries.



The generator takes in random numbers and returns an image.

This generated image is fed into the discriminator alongside a stream of images taken from the actual, ground-truth dataset

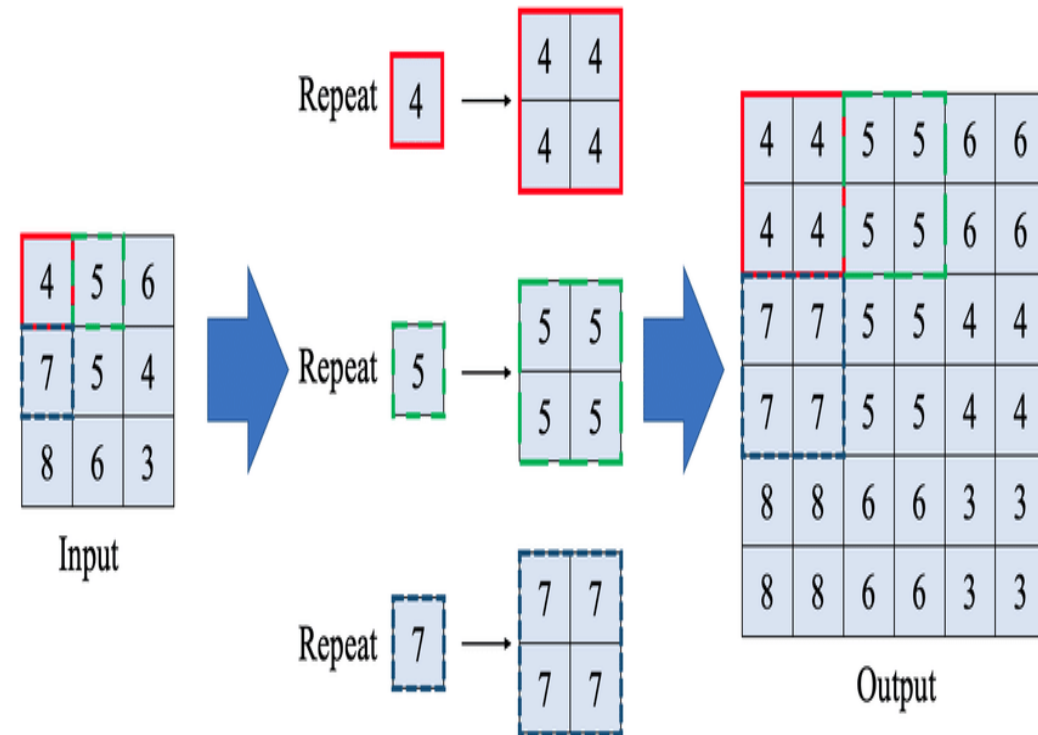

The discriminator takes in both real and fake images and returns probabilities: numbers between 0 and 1, with 1 representing a prediction of authenticity and 0 representing a prediction of fake



# Upsampling in Deep Learning


Using neural networks to generate images usually involves up-sampling from low resolution to high resolution

- Traditional convolutional neural networks use pooling layers to downsample input images.
- To scale feature maps, we use upsampling layers that scale image dimensions by repeating each row and column of input pixels.
- Keras has an upsampling layer (UpSampling2D) that scales image dimensions using a scaling factor: `keras.layers.UpSampling2D(size=(2, 2))`.
- This repeats every row and column of the image matrix two times when the scaling factor is (2, 2).
- If the scaling factor is (3, 3), the upsampling layer repeats each row and column three times.



## Conv2DTranspose in Deep Learning

- Conv2DTranspose, also known as transposed convolution or deconvolution, is used to reconstruct the spatial resolution of the input.
- It performs the reverse operation of a standard convolution.
- The process involves applying a kernel to the input data and introducing padding to expand the output size.
- Strides determine the step size for moving the kernel, with larger strides resulting in larger upsampled outputs.
- Conv2DTranspose is commonly used in image generation, semantic segmentation, and autoencoders.



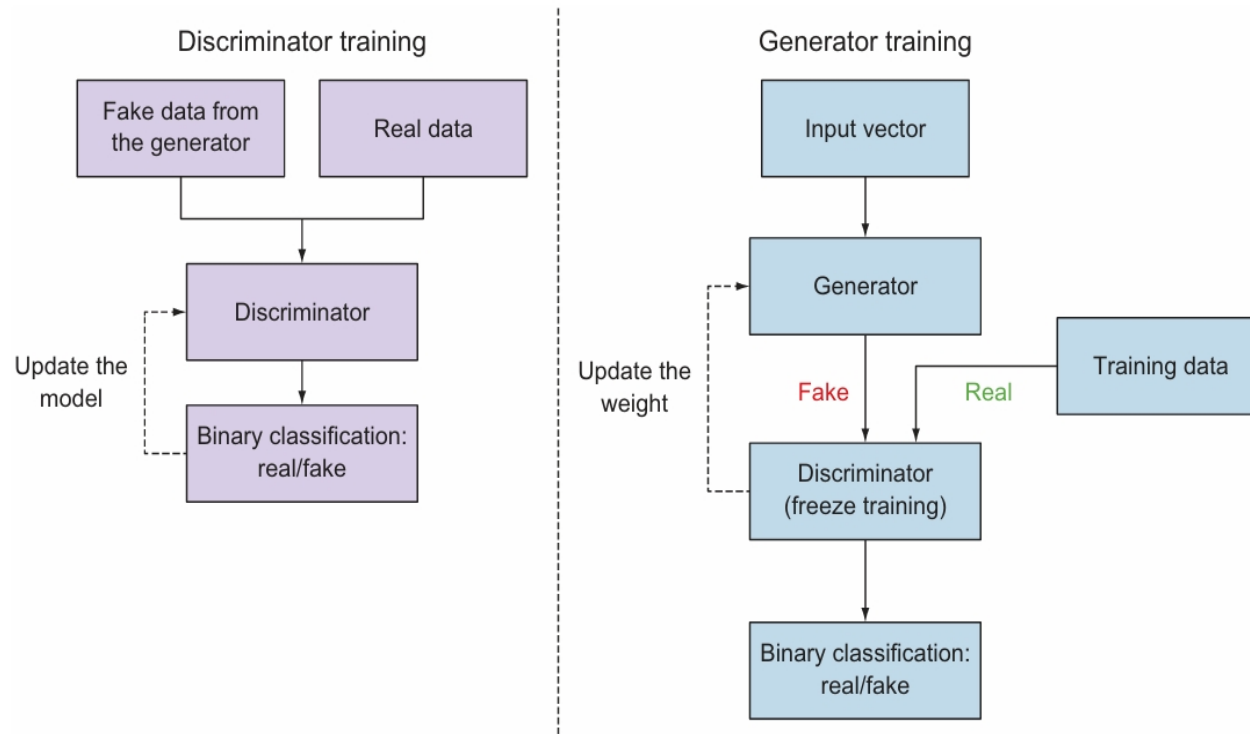


# Training the GAN

The process of training GAN models involves two processes:

1. **Train the discriminator.** This is a straightforward supervised training process. The network is given labeled images coming from the generator (fake) and the training data (real), and it learns to classify between real and fake images with a sigmoid prediction output. Nothing new here.

2. **Train the generator.** This process is a little tricky. The generator model cannot be trained alone like the discriminator. It needs the discriminator model to tell it whether it did a good job of faking images. So, we create a combined network to train the generator, composed of both discriminator and generator models.



## Training of GAN


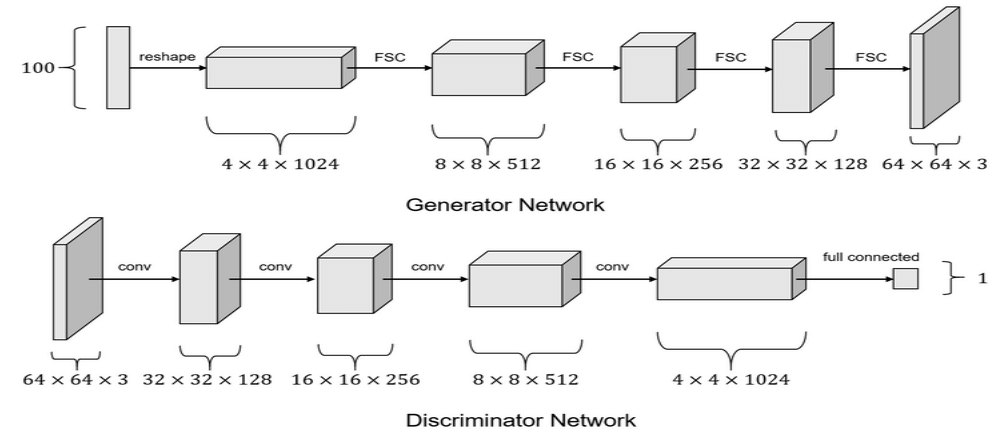




# Types of GANS



## Deep Convolutional GAN (DCGAN):

- An extension of the vanilla GAN that incorporates convolutional layers, which are particularly effective for image data.
- **Architecture:**
- Uses deep convolutional layers in the generator and discriminator instead of fully connected layers.
- Incorporates techniques such as batch normalization and leaky ReLU activations to improve training stability.
- **Advantages:**
- Improved quality of generated images.
- More stable training compared to vanilla GANs.



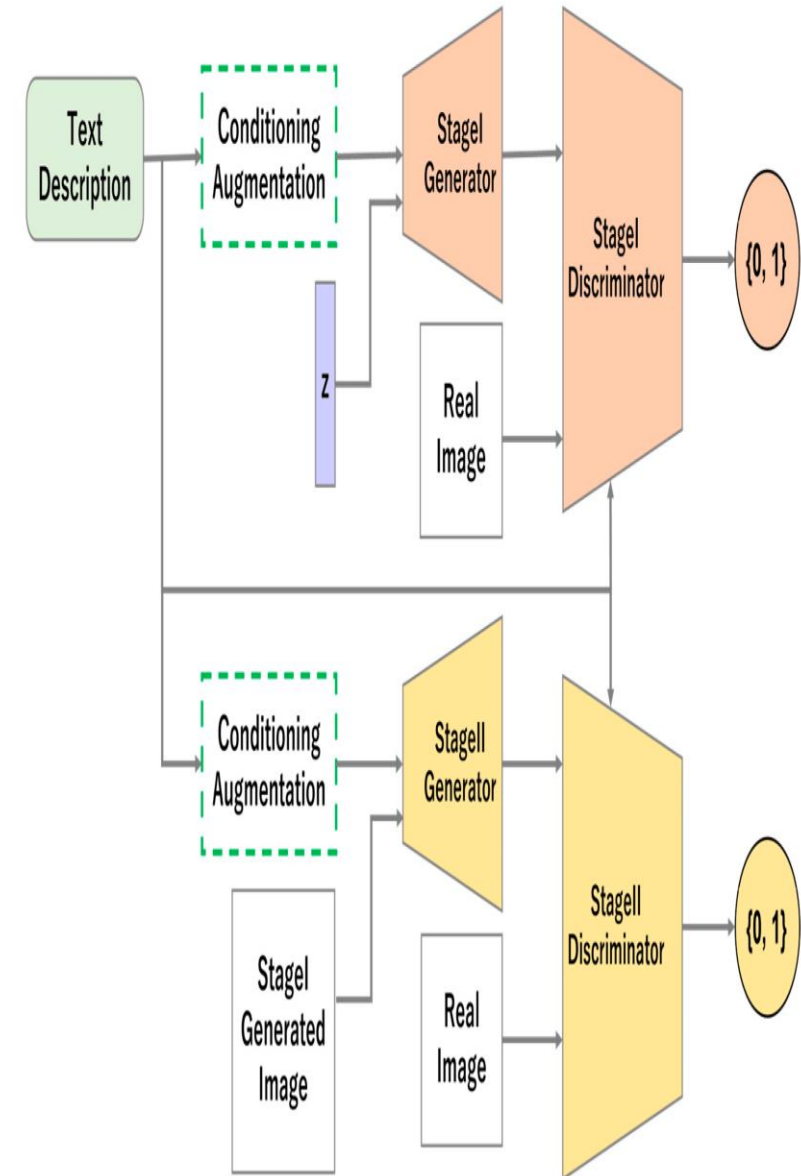
## Conditional GAN (cGAN):

- A type of GAN where both the generator and discriminator receive additional information (such as class labels) as input.
- **Architecture:**
  - The generator conditions its output on the additional information, allowing control over the type of data generated.
  - The discriminator also receives this information and uses it to determine if the generated sample matches the condition.
- **Applications:**
  - Can be used for tasks like image-to-image translation, where the model learns to generate a specific type of image based on input conditions (e.g., generating photos from sketches).




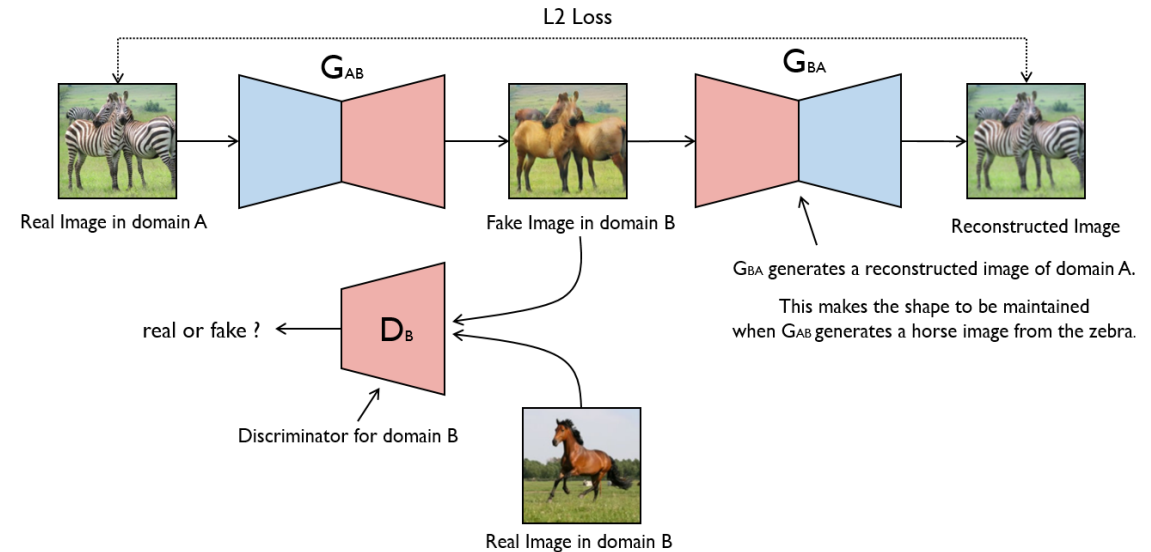
# StackGAN (Stacked Generative Adversarial Networks)

- StackGAN is a two-stage generative adversarial network (GAN) designed to produce high-resolution and photorealistic images from textual descriptions.
- Two-Stage Architecture:
  - **Stage-I GAN:**
    - Input: Text embedding, which captures the semantic meaning of the description.
    - Output: A low-resolution (64x64) image that represents the basic layout, shapes, and primary colors described by the text.
    - Purpose: Establishes a rough sketch of the final image.
  - **Stage-II GAN:**
    - Input: The low-resolution image from Stage-I and the original text embedding.
    - Output: A high-resolution (256x256) image with finer details and improved realism.
    - Purpose: Refines the initial image, adding details and enhancing the overall quality to produce a photorealistic image.
- **Applications:**
  - Creating visual content from textual descriptions in art, fashion, and design.



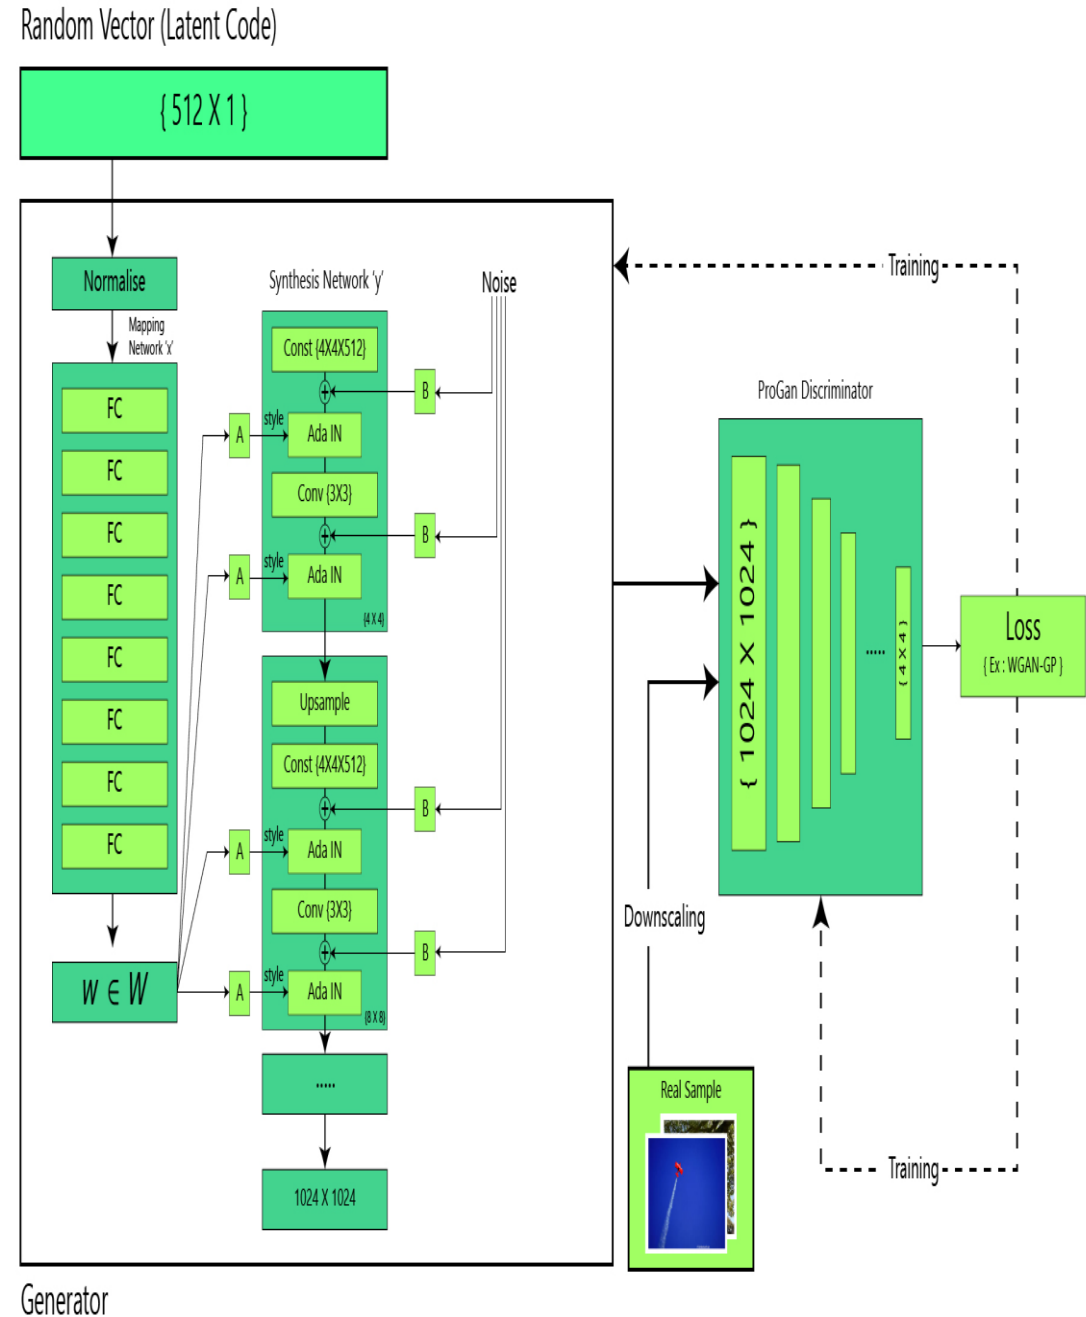
# CycleGAN:

- Designed for unpaired image-to-image translation, where the model learns to translate between two domains without requiring paired examples.
- Architecture:**
- Uses two generators and two discriminators: one set for each direction of translation (e.g., photo to painting and painting to photo).
- Employs a cycle consistency loss to ensure that translating an image to the target domain and back to the source domain results in the original image.
- Applications:**
- Useful for tasks where paired training data is not available, such as converting photos of horses to zebras and vice versa.



## StyleGAN:

- An advanced GAN architecture designed to generate high-quality images with control over the style of the generated content.
- **Architecture:**
- Introduces a style-based generator architecture, where the input latent vector controls the style of the output image at different levels of detail.
- Employs techniques like adaptive instance normalization (AdaIN) to blend different styles.
- **Advantages:**
- Capable of generating highly realistic and detailed images with controllable styles.
- Widely used for tasks like creating realistic human faces and artistic style transfer.



# Application of GANS

## 1. Image-to-Image Translation:

- **Style Transfer:**
- GANs can transfer the style of one image onto another, such as converting a photo into a painting with the style of Van Gogh or Picasso.
- CycleGAN can translate between different image domains without needing paired training examples, such as turning photos of horses into zebras.
- **Semantic Image Synthesis:**
- GANs can generate images from semantic maps or sketches.
- Pix2Pix translates sketches into realistic images, which can be used in design and artistic applications.

Day to night




Input



Output


Input





Output



Input



Output





## 2. Text-to-image Translation

- StackGAN, a GAN network designed specifically for this task, generates  $256 \times 256$  photorealistic images conditioned on text descriptions
- Synthesis of high-quality images from text descriptions is a challenging problem in CV. Samples generated by existing text-to-image approaches can roughly reflect the meaning of the given descriptions, but they fail to contain necessary details and vivid object parts.
- The GAN network that was built for this application is the stacked generative adversarial network (StackGAN). were able to generate  $256 \times 256$  photo realistic images conditioned on text descriptions.

This bird is white with some black on its head and wings, and has a long orange beak.



a) StackGAN Stage-I  
 $64 \times 64$  images



b) StackGAN Stage-II  
 $256 \times 256$  images

This bird has a yellow belly and tarsus, gray back, wings, and brown throat, nape with a black face.



This flower has overlapping pink pointed petals surrounding a ring of short yellow filaments.



### 3. Image super-resolution GAN (SRGAN)

- A certain type of GAN models can be used to convert low-resolution images into high resolution images.
- This type is called a super-resolution generative adversarial networks (SRGAN).
- SRGAN was able to create a very high-resolution image.
- **Upscaling Low-Resolution Images:**
- SRGANs excel at taking low-resolution images (e.g., from old photographs, surveillance cameras, or compressed videos) and generating high-resolution versions that retain detail and sharpness.
- **Super-resolution of Medical Scans:**
- SRGANs can be applied to medical images like MRIs or CT scans, increasing their resolution without compromising diagnostic information.




LR image

Original



4x HR image



SRGAN





## 4. Data Augmentation:

- GANs generate additional training data to improve the performance of machine learning models, particularly in scenarios with limited real data.
- GANs are used in medical imaging to augment datasets, providing more examples for training diagnostic models.
- Data augmentation is a technique used to increase the diversity and volume of training data without collecting new data. Generative Adversarial Networks (GANs) play a pivotal role in this process by generating realistic synthetic data that enhances the performance and robustness of machine learning models, especially in scenarios with limited real data.



## 5.Photo to emojis

- GANs can transform a photo of a human face into a corresponding emoji, capturing essential features and expressions.
- Demonstrates the creative potential of GANs in stylizing images for personalized digital expression.
- Generator Network: Converts a photo into an emoji.
- Discriminator Network: Ensures the generated emoji resembles real emojis by distinguishing between real and fake ones.
- **Applications:**
  - Personalized Emojis: Create custom emojis for social media and messaging apps.
  - Digital Creativity: Enhance user interaction with unique, expressive emojis.





## 6.Using GANs: Face App

- Mobile application utilizing GANs for various facial transformations such as aging, gender swapping, adding smiles, and altering hairstyles and colors.
- GANs can be used to generate images of people at different ages, allowing users to visualize how they might look in the future or to see what they might have looked like in the past..
- **Aging Effects:** Simulate how a person might look as they age.
- **Gender Swap:** Transform a person's facial features to appear as the opposite gender.
- **Smile Addition:** Add realistic smiles to photos where the person is not smiling.
- **Hair Style and Color Changes:** Alter hairstyles and colors realistically.
- **Generator Network:** Applies the desired transformation to the input photo.
- **Discriminator Network:** Ensures the transformed photo looks realistic and authentic.
- **Adversarial Training:** Continual improvement of transformations through feedback between generator and discriminator.
- Social Media: Share fun and engaging photos with friends and followers.


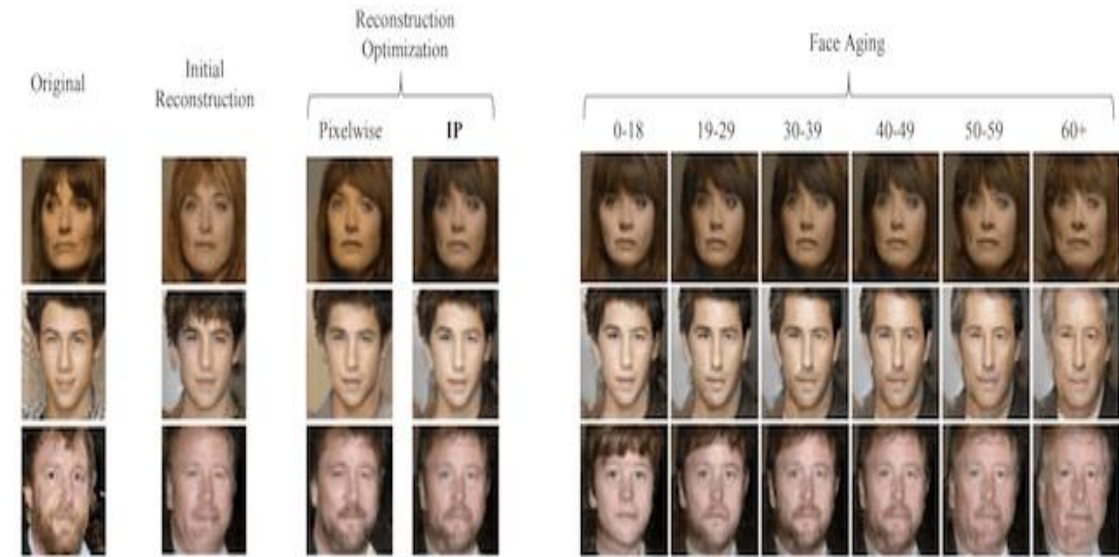


Figure 1. Demonstration of our aging simulation results (image:

## Other Applications

- **Image Generation and Enhancement:**
- **Image Inpainting:** GANs can fill in missing parts of an image, useful in photo restoration and editing.
- **Video Generation and Editing:**
- **Video Prediction:** GANs can predict future frames in a video, useful for video compression and enhancement.
- **Deepfake Technology:** GANs can create realistic videos of people saying or doing things they never actually did, with applications in entertainment and, unfortunately, potential misuse in misinformation.
- **Art and Design:**
- **Art Creation:** GANs can generate new pieces of art, providing tools for artists and designers to explore new creative boundaries.
- **Style Transfer:** GANs can apply the style of one image (e.g., a painting) to another image (e.g., a photograph), useful in creative and commercial applications.

- **Speech and Audio Processing:**
- **Voice Synthesis:** GANs can generate realistic human speech, useful in virtual assistants and automated customer service.
- **Music Generation:** GANs can create new music tracks, assisting musicians in the composition process.
- **Fashion and Retail:**
- **Virtual Try-On:** GANs can generate realistic images of how clothes will look on different body types, aiding online shopping experiences.
- **Design Prototyping:** GANs can create new fashion designs, helping designers in the ideation and prototyping stages.
- **Language Learning:**
- **Language Practice:** GANs can generate conversational agents that provide students with realistic practice in foreign languages, improving their speaking and listening skills.
- **Content Translation:** GANs can assist in translating educational materials into multiple languages, making education more accessible to non-native speakers.

- **Content Creation and Personalization:**
- **Automated Content Generation:** GANs can generate educational content such as practice problems, quizzes, and even explanatory text, tailored to different learning levels and styles.
- **Personalized Learning Materials:** GANs can create customized textbooks and learning resources that cater to the specific needs and learning pace of individual students.

Thank You